

# Machine learning Homework- Linear regression

Abinav Ravi Venkatakrishnan - 03694216 and Abhijeet Parida - 03679676

November 11, 2018

## 1 Least Square Regression

### Problem 2

**Given:** The error function

$$E_{weighted}(w) = \frac{1}{2} \sum_{i=1}^N t_i [w^T \phi(x_i) - y_i]^2 \quad (1)$$

We know that taking the first derivative of the cost function gives the minimum of the error function.

$$\nabla E_{weighted}(w) = \nabla_w \frac{1}{2} \sum_{i=1}^N t_i [w^T \phi(x_i) - y_i]^2 = 0 \quad (2)$$

Now the RHS can be written as a matrix equation with T as the diagonal matrix.

$$\nabla_w E_{weighted} = \nabla_w \frac{1}{2} [\Phi^T \mathbf{T} w - y]^T [\Phi^T \mathbf{T} w - y] \quad (3)$$

If  $\mathbf{T} = \mathbf{I}$  (Identity Matrix) hence giving the solution

$$w_* = (\Phi^T \Phi)^{-1} \Phi^T y \quad (4)$$

$t_i$  acts as a multiplication factor for variance in  $y_i$  due to  $x_i$

when the data points for which there are exact copies in the dataset. The error contribution due to these repeated points are increased.

## 2 Ridge regression

### Problem 3

Before Augmentation,

$$E_{ridge} = \frac{1}{2} \sum_{i=1}^N (w_N^T \phi(x_i) - y_i)^2 + \frac{\lambda}{2} w_N^T w_N \quad (5)$$

After Augmentation,

$$E_{ridge} = \frac{1}{2} \sum_{i=1}^{N+M} (w_{N+M}^T \phi(x_i) - y_i)^2 + \frac{\lambda}{2} w_{N+M}^T w_{N+M}$$

This can be split as,

$$E_{ridge} = \frac{1}{2} \sum_{i=1}^{N+M} (w_{N+M}^T \phi(x_i) - y_i)^2 + \frac{\lambda}{2} w_{N+M}^T w_{N+M} \quad (6)$$

$$= \frac{1}{2} \sum_{i=1}^N (w_N^T \phi(x_i) - y_i)^2 + \frac{\lambda}{2} w_N^T w_N + \frac{1}{2} \sum_{i=N+1}^{N+M} (w_{N+M}^T \phi(x_i) - y_i)^2 + \frac{\lambda}{2} w_{N+M}^T w_{N+M} \quad (7)$$

$$(8)$$

In 8, the term  $w_{N+M} = 0$ , because

$$w_{N+M} = (X_{N+M}^T X_{N+M})^{-1} X_{N+M}^T y_{N+M}$$

because

$$y_{N+M} = 0$$

Therefore 8 is equal to 5

### 3 Bayesian linear regression

#### Problem 4

We know that

posterior  $\propto$  likelihood  $\times$  prior

we will work in the log scale to easily deal with the  $\prod$

$$\begin{aligned} \log(\mathcal{N}(\mathbf{w}|\mathbf{m}_N, \beta^{-1}\mathbf{S}_N)) + \log(\text{Gamma}(\beta|a_N, b_N)) = & \sum_{i=1}^N \log(\mathcal{N}(y_i|\mathbf{w}^T\phi(x_i), \beta^{-1})) \\ & + \log(\mathcal{N}(\mathbf{w}|\mathbf{m}_0, \beta^{-1}\mathbf{S}_0)) \\ & + \log(\text{Gamma}(\beta|a_0, b_0)) \end{aligned}$$

By simple pattern matching after resolution we can easily find the unknowns and will hence prove the validity of the claim.

# homework\_03\_notebook

November 9, 2018

## 1 Programming assignment 2: Linear regression

```
In [1]: import numpy as np
```

```
from sklearn.datasets import load_boston
from sklearn.model_selection import train_test_split
```

### 1.1 Your task

In this notebook code skeleton for performing linear regression is given. Your task is to complete the functions where required. You are only allowed to use built-in Python functions, as well as any numpy functions. No other libraries / imports are allowed.

### 1.2 Load and preprocess the data

In this assignment we will work with the Boston Housing Dataset. The data consists of 506 samples. Each sample represents a district in the city of Boston and has 13 features, such as crime rate or taxation level. The regression target is the median house price in the given district (in \$1000's).

More details can be found here: <http://lib.stat.cmu.edu/datasets/boston>

```
In [2]: X , y = load_boston(return_X_y=True)
```

```
# Add a vector of ones to the data matrix to absorb the bias term
# (Recall slide #7 from the lecture)
X = np.hstack([np.ones([X.shape[0], 1]), X])
# From now on, D refers to the number of features in the AUGMENTED dataset
# (i.e. including the dummy '1' feature for the absorbed bias term)

# Split into train and test
test_size = 0.2
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size)
```

### 1.3 Task 1: Fit standard linear regression

```
In [3]: def fit_least_squares(X, y):
        """Fit ordinary least squares model to the data.
        Parameters
```

```

-----
X : array, shape [N, D]
    (Augmented) feature matrix.
y : array, shape [N]
    Regression targets.

Returns
-----
w : array, shape [D]
    Optimal regression coefficients (w[0] is the bias term).

"""
# TODO
w = np.zeros_like(X.shape[1])
X_transpose = X.transpose((1,0))
X_star = X_transpose.dot(X)
X_inverse = np.linalg.inv(X_star)
w = X_inverse.dot(X_transpose).dot(y)
return w

```

## 1.4 Task 2: Fit ridge regression

```

In [10]: def fit_ridge(X, y, reg_strength):
    """Fit ridge regression model to the data.

    Parameters
    -----
    X : array, shape [N, D]
        (Augmented) feature matrix.
    y : array, shape [N]
        Regression targets.
    reg_strength : float
        L2 regularization strength (denoted by lambda in the lecture)

    Returns
    -----
    w : array, shape [D]
        Optimal regression coefficients (w[0] is the bias term).

    """
    # TODO
    w = np.zeros_like(X.shape[1])
    X_transpose = X.transpose((1,0))
    X_star = X_transpose.dot(X) + reg_strength*np.eye(X.shape[1])
    X_inverse = np.linalg.inv(X_star)
    w = X_inverse.dot(X_transpose).dot(y)
    return w

```

## 1.5 Task 3: Generate predictions for new data

```
In [11]: def predict_linear_model(X, w):  
    """Generate predictions for the given samples.  
  
    Parameters  
    -----  
    X : array, shape [N, D]  
        (Augmented) feature matrix.  
    w : array, shape [D]  
        Regression coefficients.  
  
    Returns  
    -----  
    y_pred : array, shape [N]  
        Predicted regression targets for the input data.  
  
    """  
    # TODO  
    y_pred = np.dot(X,w)  
    return y_pred
```

## 1.6 Task 4: Mean squared error

```
In [6]: def mean_squared_error(y_true, y_pred):  
    """Compute mean squared error between true and predicted regression targets.  
  
    Reference: `https://en.wikipedia.org/wiki/Mean_squared_error`  
  
    Parameters  
    -----  
    y_true : array  
        True regression targets.  
    y_pred : array  
        Predicted regression targets.  
  
    Returns  
    -----  
    mse : float  
        Mean squared error.  
  
    """  
    # TODO  
    mse = np.mean((y_pred-y_true)**2)  
    return mse
```

## 1.7 Compare the two models

The reference implementation produces \* MSE for Least squares  $\approx 23.98$  \* MSE for Ridge regression  $\approx 21.05$

Your results might be slightly (i.e.  $\pm 1\%$ ) different from the reference solution due to numerical reasons.

```
In [12]: # Load the data
np.random.seed(1234)
X, y = load_boston(return_X_y=True)
X = np.hstack([np.ones([X.shape[0], 1]), X])
test_size = 0.2
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size)

# Ordinary least squares regression
w_ls = fit_least_squares(X_train, y_train)
y_pred_ls = predict_linear_model(X_test, w_ls)
mse_ls = mean_squared_error(y_test, y_pred_ls)
print('MSE for Least squares = {}'.format(mse_ls))

# Ridge regression
reg_strength = 1
w_ridge = fit_ridge(X_train, y_train, reg_strength)
y_pred_ridge = predict_linear_model(X_test, w_ridge)
mse_ridge = mean_squared_error(y_test, y_pred_ridge)
print('MSE for Ridge regression = {}'.format(mse_ridge))

MSE for Least squares = 23.984307611781773
MSE for Ridge regression = 21.051487033772275
```