

# Machine learning Homework- Optimisation

Abinav Ravi Venkatakrishnan - 03694216 and Abhijeet Parida - 03679676

November 24, 2018

## Problem 1:

1. we know that sum of convex function is a convex function.  $x^2, 2y, \cos(\sin(\sqrt{\pi}))$  are convex functions because linear and constant are convex.  
 $-\min\{-x^2, \log(y)\} = \max\{x^2, -\log(y)\}$ . The max is convex if both functions are convex which is true for our case. So, it is a convex function.
2. we know that the point-wise addition of two convex function is convex; but here  $\log(x)$  and  $-x^3$  are concave; so overall is concave.
3.  $-\min\{\log(3x+1), -x^4 - 3x^2 + 8x - 42\} = \max\{-\log(3x+1), x^4 + 3x^2 - 8x + 42\}$  The max is convex if both functions are convex which is true for our case. So, it is a convex function.
4. by plotting Fig. 1, we see that the function  $-x^2y$  is not convex. So, it is not a convex function.




Figure 1:

## Problem 2:

we know that the functions  $f_1$  and  $f_2$  are convex in  $\mathcal{R}$ .

so,

$$\lambda_1 f_1(x_1) + (1 - \lambda_1) f_1(x_2) \geq f_1(\lambda_1 x_1 + (1 - \lambda_1) x_2) \quad (1)$$

and

$$\lambda_1 f_2(x_1) + (1 - \lambda_1) f_2(x_2) \geq f_2(\lambda_1 x_1 + (1 - \lambda_1) x_2) \quad (2)$$

$h(x) = \max(f_1(x), f_2(x))$  and given  $f_1$  and  $f_2$  are convex can be written as

$$\begin{cases} f_1(x) & ; x \geq C \\ f_2(x) & ; x \leq C \end{cases}$$

for  $x_1, x_2 \geq C$ ; Eqn 1 is true and for  $x_1, x_2 \leq C$ ; Eqn 2 is true.

for the case  $x_2 \geq C$  and  $x_1 \leq C$ . We have to show two scenario to be true.

$$\lambda_1 f_1(x_1) + (1 - \lambda_1) f_2(x_2) \geq f_1(\lambda_1 x_1 + (1 - \lambda_1) x_2) \quad (3)$$

Eq. 3 is true because Eq. 1 is true and on the LHS  $(1 - \lambda_1) f_2(x_2) > (1 - \lambda_1) f_1(x_2)$  from definition of  $h(x)$ .

$$\lambda_1 f_1(x_1) + (1 - \lambda_1) f_2(x_2) \geq f_2(\lambda_1 x_1 + (1 - \lambda_1) x_2) \quad (4)$$

Eq. 4 is true because Eq. 2 is true and on the LHS  $\lambda_1 f_1(x_1) > \lambda_1 f_2(x_1)$  from definition of  $h(x)$ .

### Problem 3:

we know that  $f_1$  and  $f_2$  are convex and their exist a minima  $x_1$  and  $x_2$  such that

$$f'_1(x_1) = 0 ; f''_1(x_1) > 0 \quad (5)$$

$$f'_2(x_2) = 0 ; f''_2(x_2) > 0 \quad (6)$$

for our function  $f_1(f_2(x))$ , the local minima can be found at

$$f'_1(f_2(x)) \cdot f'_2(x) = 0 \quad (7)$$

From Eq. 5 and Eq. 6 the possible candidates for minima of  $f_1 f_2(x)$  is  $x = x_2$  and  $f_2(x) = x_1$ .

For the critical point  $x = x_2$  is minima if  $(f_1 f_2(x))'' > 0$

$$f''_2(x) \cdot f'_1(f_2(x)) + f'_2(x) \cdot f''_1(f_2(x)) \cdot f'_2(x) \quad (8)$$

we know that  $f'_2(x_2) \cdot f''_1(f_2(x_2)) \cdot f'_2(x_2) = 0$  from Eq. 6 and in the term  $f''_2(x) \cdot f'_1(f_2(x)) \cdot f'_2(x)$  is +ve from Eq. 6. Therefore, the final sign of the Eq.8 depends on the sign of  $f'_1(f_2(x)) \cdot f'_2(x)$ . Since it is positive/negative we cannot say it is convex always.

For the critical point  $f_2(x) = x_1$  is minima if  $(f_1 f_2(x))'' > 0$ .

we know that  $f''_2(x) \cdot f'_1(f_2(x)) = 0$  from Eq. 7 and in the term  $f'_2(x) \cdot f''_1(f_2(x)) \cdot f'_2(x)$  is always +ve from Eq. 6. Therefore, the statement is true in this case.

Overall, the statement is true for some  $f_1, f_2$  and not true for other  $f_2, f_1$

### Problem 4:

Theorem: Suppose  $f : x \mapsto \mathcal{R}$  is a differentiable function and  $\mathcal{R}$  is convex. Then  $f$  is convex iff for  $x, y \in \mathcal{X}$  then

$$f(y) \geq f(x) + (y - x)^T \nabla f(x) \quad (9)$$

for  $\nabla f(x) = 0$ , we get  $f(y) \geq f(x)$  hence  $f(x)$  is the lowest possible value when  $\nabla f(x) = 0$

# Programming assignment 5: Optimization: Logistic regression

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score
```

## Your task

In this notebook code skeleton for performing logistic regression with gradient descent is given. Your task is to complete the functions where required. You are only allowed to use built-in Python functions, as well as any `numpy` functions. No other libraries / imports are allowed.

For numerical reasons, we actually minimize the following loss function

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} NLL(\mathbf{w}) + \frac{1}{2} \lambda \|\mathbf{w}\|_2^2$$

where  $NLL(\mathbf{w})$  is the negative log-likelihood function, as defined in the lecture (Eq. 33)

## Exporting the results to PDF

Once you complete the assignments, export the entire notebook as PDF and attach it to your homework solutions. The best way of doing that is

1. Run all the cells of the notebook.
2. Download the notebook in HTML (click File > Download as > .html)
3. Convert the HTML to PDF using e.g. <https://www.sejda.com/html-to-pdf> or `wkhtmltopdf` for Linux ([tutorial](#))
4. Concatenate your solutions for other tasks with the output of Step 3. On a Linux machine you can simply use `pdfunite`, there are similar tools for other platforms too. You can only upload a single PDF file to Moodle.

This way is preferred to using `nbconvert`, since `nbconvert` clips lines that exceed page width and makes your code harder to grade.

## Load and preprocess the data

In this assignment we will work with the UCI ML Breast Cancer Wisconsin (Diagnostic) dataset <https://goo.gl/U2Uwz2>.

Features are computed from a digitized image of a fine needle aspirate (FNA) of a breast mass. They describe characteristics of the cell nuclei present in the image. There are 212 malignant examples and 357 benign examples.

```
In [2]:
```

```

X, y = load_breast_cancer(return_X_y=True)

# Add a vector of ones to the data matrix to absorb the bias term
X = np.hstack([np.ones([X.shape[0], 1]), X])

# Set the random seed so that we have reproducible experiments
np.random.seed(123)

# Split into train and test
test_size = 0.3
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size)

```

## Task 1: Implement the sigmoid function

```

In [3]: def sigmoid(t):
        """
        Applies the sigmoid function elementwise to the input data.

        Parameters
        -----
        t : array, arbitrary shape
            Input data.

        Returns
        -----
        t_sigmoid : array, arbitrary shape.
            Data after applying the sigmoid function.
        """
        # TODO
        sigmoid = 1/(1+np.exp(-t))
        return sigmoid

```

## Task 2: Implement the negative log likelihood

As defined in Eq. 33

```

In [4]: def negative_log_likelihood(X, y, w):
        """
        Negative Log Likelihood of the Logistic Regression.

        Parameters
        -----
        X : array, shape [N, D]
            (Augmented) feature matrix.
        y : array, shape [N]
            Classification targets.
        w : array, shape [D]
            Regression coefficients (w[0] is the bias term).

        Returns
        -----
        nll : float
            The negative log likelihood.
        """
        # TODO
        value = sigmoid(np.dot(X,w))
        nll = np.sum((y*np.log(value))+(1-y)*(np.log(1-value)))

        return (nll)

```

## Computing the loss function $\mathcal{L}(\mathbf{w})$ (nothing to do here)

```
In [5]: def compute_loss(X, y, w, lambda):  
    """  
    Negative Log Likelihood of the Logistic Regression.  
  
    Parameters  
    -----  
    X : array, shape [N, D]  
        (Augmented) feature matrix.  
    y : array, shape [N]  
        Classification targets.  
    w : array, shape [D]  
        Regression coefficients (w[0] is the bias term).  
    lambda : float  
        L2 regularization strength.  
  
    Returns  
    -----  
    loss : float  
        Loss of the regularized logistic regression model.  
    """  
    # The bias term w[0] is not regularized by convention  
    return negative_log_likelihood(X, y, w) / len(y) + lambda * np.linalg.norm(w  
[1:])**2
```

## Task 3: Implement the gradient $\nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w})$

Make sure that you compute the gradient of the loss function  $\mathcal{L}(\mathbf{w})$  (not simply the NLL!)

```
In [6]: def get_gradient(X, y, w, mini_batch_indices, lambda):  
    """  
    Calculates the gradient (full or mini-batch) of the negative log likelihood  
    d w.r.t. w.  
  
    Parameters  
    -----  
    X : array, shape [N, D]  
        (Augmented) feature matrix.  
    y : array, shape [N]  
        Classification targets.  
    w : array, shape [D]  
        Regression coefficients (w[0] is the bias term).  
    mini_batch_indices: array, shape [mini_batch_size]  
        The indices of the data points to be included in the (stochastic) calculation  
        of the gradient.  
        This includes the full batch gradient as well, if mini_batch_indices =  
        np.arange(n_train).  
    lambda: float  
        Regularization strength. lambda = 0 means having no regularization.  
  
    Returns  
    -----  
    dw : array, shape [D]  
        Gradient w.r.t. w.  
    """  
    # TODO  
    n_batch = mini_batch_indices.shape[0]  
    nll_gradient = np.dot(X[mini_batch_indices].T, sigmoid(np.dot(X[mini_batch_i  
ndices], w)) - y[mini_batch_indices])  
    ones = np.ones(w.shape)
```

```

ones[0] = 0
reg_gradient = lambda * ones * w
grad = nll_gradient / n_batch + reg_gradient
return grad

```

## Train the logistic regression model (nothing to do here)

```

In [7]: def logistic_regression(X, y, num_steps, learning_rate, mini_batch_size, lambda,
        verbose):
    """
    Performs logistic regression with (stochastic) gradient descent.

    Parameters
    -----
    X : array, shape [N, D]
        (Augmented) feature matrix.
    y : array, shape [N]
        Classification targets.
    num_steps : int
        Number of steps of gradient descent to perform.
    learning_rate: float
        The learning rate to use when updating the parameters w.
    mini_batch_size: int
        The number of examples in each mini-batch.
        If mini_batch_size=n_train we perform full batch gradient descent.
    lambda: float
        Regularization strength. lambda = 0 means having no regularization.
    verbose : bool
        Whether to print the loss during optimization.

    Returns
    -----
    w : array, shape [D]
        Optimal regression coefficients (w[0] is the bias term).
    trace: list
        Trace of the loss function after each step of gradient descent.
    """

    trace = [] # saves the value of loss every 50 iterations to be able to plot
    it later
    n_train = X.shape[0] # number of training instances

    w = np.zeros(X.shape[1]) # initialize the parameters to zeros

    # run gradient descent for a given number of steps
    for step in range(num_steps):
        permuted_idx = np.random.permutation(n_train) # shuffle the data

        # go over each mini-batch and update the parameters
        # if mini_batch_size = n_train we perform full batch GD and this loop runs
        # only once
        for idx in range(0, n_train, mini_batch_size):
            # get the random indices to be included in the mini batch
            mini_batch_indices = permuted_idx[idx:idx+mini_batch_size]
            gradient = get_gradient(X, y, w, mini_batch_indices, lambda)

            # update the parameters
            w = w - learning_rate * gradient

        # calculate and save the current loss value every 50 iterations
        if step % 50 == 0:
            loss = compute_loss(X, y, w, lambda)
            trace.append(loss)

```

```

        # print loss to monitor the progress
        if verbose:
            print('Step {0}, loss = {1:.4f}'.format(step, loss))
    return w, trace

```

## Task 4: Implement the function to obtain the predictions

```

In [8]: def predict(X, w):
        """
        Parameters
        -----
        X : array, shape [N_test, D]
            (Augmented) feature matrix.
        w : array, shape [D]
            Regression coefficients (w[0] is the bias term).

        Returns
        -----
        y_pred : array, shape [N_test]
            A binary array of predictions.
        """
        # TODO
        return (sigmoid(np.dot(X, w)) > 0.5).astype(np.int)

```

## Full batch gradient descent

```

In [9]: # Change this to True if you want to see loss values over iterations.
        verbose = True

```

```

In [10]: n_train = X_train.shape[0]
        w_full, trace_full = logistic_regression(X_train,
                                                y_train,
                                                num_steps=8000,
                                                learning_rate=1e-5,
                                                mini_batch_size=n_train,
                                                lmbda=0.1,
                                                verbose=verbose)

```

```

Step 0, loss = -0.7427
Step 50, loss = -0.9390
Step 100, loss = -0.5167
Step 150, loss = -0.3868
Step 200, loss = -0.3675
Step 250, loss = -0.3522
Step 300, loss = -0.3395
Step 350, loss = -0.3288
Step 400, loss = -0.3197
Step 450, loss = -0.3117
Step 500, loss = -0.3048
Step 550, loss = -0.2986
Step 600, loss = -0.2931
Step 650, loss = -0.2882
Step 700, loss = -0.2837
Step 750, loss = -0.2796
Step 800, loss = -0.2759
Step 850, loss = -0.2725
Step 900, loss = -0.2694

Step 950, loss = -0.2665
Step 1000, loss = -0.2639
Step 1050, loss = -0.2614
Step 1100, loss = -0.2591

```

Step 1150, loss = -0.2569  
Step 1200, loss = -0.2549  
Step 1250, loss = -0.2530  
Step 1300, loss = -0.2513  
Step 1350, loss = -0.2496  
Step 1400, loss = -0.2481  
Step 1450, loss = -0.2466  
Step 1500, loss = -0.2452  
Step 1550, loss = -0.2439  
Step 1600, loss = -0.2427  
Step 1650, loss = -0.2415  
Step 1700, loss = -0.2404  
Step 1750, loss = -0.2393  
Step 1800, loss = -0.2383  
Step 1850, loss = -0.2373  
Step 1900, loss = -0.2364  
Step 1950, loss = -0.2356  
Step 2000, loss = -0.2347  
Step 2050, loss = -0.2339  
Step 2100, loss = -0.2332  
Step 2150, loss = -0.2325  
Step 2200, loss = -0.2318  
Step 2250, loss = -0.2311  
Step 2300, loss = -0.2305  
Step 2350, loss = -0.2299  
Step 2400, loss = -0.2293  
Step 2450, loss = -0.2287  
Step 2500, loss = -0.2282  
Step 2550, loss = -0.2276  
Step 2600, loss = -0.2271  
Step 2650, loss = -0.2266  
Step 2700, loss = -0.2262  
Step 2750, loss = -0.2257  
Step 2800, loss = -0.2253  
Step 2850, loss = -0.2249  
Step 2900, loss = -0.2245  
Step 2950, loss = -0.2241  
Step 3000, loss = -0.2237  
Step 3050, loss = -0.2233  
Step 3100, loss = -0.2230  
Step 3150, loss = -0.2226  
Step 3200, loss = -0.2223  
Step 3250, loss = -0.2219  
Step 3300, loss = -0.2216  
Step 3350, loss = -0.2213  
Step 3400, loss = -0.2210  
Step 3450, loss = -0.2207  
Step 3500, loss = -0.2205  
Step 3550, loss = -0.2202  
Step 3600, loss = -0.2199  
Step 3650, loss = -0.2196  
Step 3700, loss = -0.2194  
Step 3750, loss = -0.2191  
Step 3800, loss = -0.2189  
Step 3850, loss = -0.2187  
Step 3900, loss = -0.2184  
Step 3950, loss = -0.2182  
Step 4000, loss = -0.2180  
Step 4050, loss = -0.2178  
Step 4100, loss = -0.2176  
Step 4150, loss = -0.2174  
Step 4200, loss = -0.2172  
Step 4250, loss = -0.2170  
Step 4300, loss = -0.2168  
Step 4350, loss = -0.2166  
Step 4400, loss = -0.2164



Step 4450, loss = -0.2163  
Step 4500, loss = -0.2161  
Step 4550, loss = -0.2159  
Step 4600, loss = -0.2157  
Step 4650, loss = -0.2156  
Step 4700, loss = -0.2154  
Step 4750, loss = -0.2153  
Step 4800, loss = -0.2151  
Step 4850, loss = -0.2150  
Step 4900, loss = -0.2148  
Step 4950, loss = -0.2147  
Step 5000, loss = -0.2145  
Step 5050, loss = -0.2144  
Step 5100, loss = -0.2142  
Step 5150, loss = -0.2141  
Step 5200, loss = -0.2140  
Step 5250, loss = -0.2138  
Step 5300, loss = -0.2137  
Step 5350, loss = -0.2136  
Step 5400, loss = -0.2135  
Step 5450, loss = -0.2133  
Step 5500, loss = -0.2132  
Step 5550, loss = -0.2131  
Step 5600, loss = -0.2130  
Step 5650, loss = -0.2129  
Step 5700, loss = -0.2128  
Step 5750, loss = -0.2126  
Step 5800, loss = -0.2125  
Step 5850, loss = -0.2124  
Step 5900, loss = -0.2123  
Step 5950, loss = -0.2122  
Step 6000, loss = -0.2121  
Step 6050, loss = -0.2120  
Step 6100, loss = -0.2119  
Step 6150, loss = -0.2118  
Step 6200, loss = -0.2117  
Step 6250, loss = -0.2116  
Step 6300, loss = -0.2115  
Step 6350, loss = -0.2114  
Step 6400, loss = -0.2113  
Step 6450, loss = -0.2112  
Step 6500, loss = -0.2111  
Step 6550, loss = -0.2110  
Step 6600, loss = -0.2110  
Step 6650, loss = -0.2109  
Step 6700, loss = -0.2108  
Step 6750, loss = -0.2107  
Step 6800, loss = -0.2106  
Step 6850, loss = -0.2105  
Step 6900, loss = -0.2104  
Step 6950, loss = -0.2104  
Step 7000, loss = -0.2103  
Step 7050, loss = -0.2102  
Step 7100, loss = -0.2101  
Step 7150, loss = -0.2100  
Step 7200, loss = -0.2100  
Step 7250, loss = -0.2099  
Step 7300, loss = -0.2098  
Step 7350, loss = -0.2097  
Step 7400, loss = -0.2097  
Step 7450, loss = -0.2096  
Step 7500, loss = -0.2095  
Step 7550, loss = -0.2094  
Step 7600, loss = -0.2094  
Step 7650, loss = -0.2093  
Step 7700, loss = -0.2092

```
Step 7750, loss = -0.2091
Step 7800, loss = -0.2091
Step 7850, loss = -0.2090
Step 7900, loss = -0.2089
Step 7950, loss = -0.2089
```

```
In [11]: n_train = X_train.shape[0]
w_minibatch, trace_minibatch = logistic_regression(X_train,
                                                    y_train,
                                                    num_steps=8000,
                                                    learning_rate=1e-5,
                                                    mini_batch_size=50,
                                                    lmbda=0.1,
                                                    verbose=verbose)
```

```
Step 0, loss = -1.3392
Step 50, loss = -0.3212
Step 100, loss = -0.2856
Step 150, loss = -0.2550
Step 200, loss = -0.2577
Step 250, loss = -0.2400
Step 300, loss = -0.2279
Step 350, loss = -0.2263
Step 400, loss = -0.2212
Step 450, loss = -0.2227
Step 500, loss = -0.2210
Step 550, loss = -0.2298
Step 600, loss = -0.2145
Step 650, loss = -0.2155
Step 700, loss = -0.2132
Step 750, loss = -0.2167
Step 800, loss = -0.2108
Step 850, loss = -0.2328
Step 900, loss = -0.2096
Step 950, loss = -0.2125
Step 1000, loss = -0.2090
Step 1050, loss = -0.2148
Step 1100, loss = -0.2078
Step 1150, loss = -0.2070
Step 1200, loss = -0.2074
Step 1250, loss = -0.2097
Step 1300, loss = -0.2084
Step 1350, loss = -0.2054
Step 1400, loss = -0.2060
Step 1450, loss = -0.2047
Step 1500, loss = -0.2043
Step 1550, loss = -0.2074
Step 1600, loss = -0.2148
Step 1650, loss = -0.2046
Step 1700, loss = -0.2115
Step 1750, loss = -0.2051
Step 1800, loss = -0.2030
Step 1850, loss = -0.2080
Step 1900, loss = -0.2019
Step 1950, loss = -0.2228
Step 2000, loss = -0.2040
Step 2050, loss = -0.2160
Step 2100, loss = -0.2008
```

```
Step 2150, loss = -0.2031
Step 2200, loss = -0.2053
Step 2250, loss = -0.2059
Step 2300, loss = -0.1998
Step 2350, loss = -0.2019
Step 2400, loss = -0.2105
Step 2450, loss = -0.2038
```

Step 2500, loss = -0.1994  
Step 2550, loss = -0.1990  
Step 2600, loss = -0.2104  
Step 2650, loss = -0.2001  
Step 2700, loss = -0.2027  
Step 2750, loss = -0.2041  
Step 2800, loss = -0.1976  
Step 2850, loss = -0.2050  
Step 2900, loss = -0.1973  
Step 2950, loss = -0.2133  
Step 3000, loss = -0.1973  
Step 3050, loss = -0.1973  
Step 3100, loss = -0.1996  
Step 3150, loss = -0.2095  
Step 3200, loss = -0.2022  
Step 3250, loss = -0.1960  
Step 3300, loss = -0.1958  
Step 3350, loss = -0.1998  
Step 3400, loss = -0.1969  
Step 3450, loss = -0.2010  
Step 3500, loss = -0.2050  
Step 3550, loss = -0.2082  
Step 3600, loss = -0.1948  
Step 3650, loss = -0.2173  
Step 3700, loss = -0.1959  
Step 3750, loss = -0.1957  
Step 3800, loss = -0.1953  
Step 3850, loss = -0.1993  
Step 3900, loss = -0.1961  
Step 3950, loss = -0.2125  
Step 4000, loss = -0.1988  
Step 4050, loss = -0.1935  
Step 4100, loss = -0.2007  
Step 4150, loss = -0.1936  
Step 4200, loss = -0.1950  
Step 4250, loss = -0.2072  
Step 4300, loss = -0.1935  
Step 4350, loss = -0.1927  
Step 4400, loss = -0.2200  
Step 4450, loss = -0.1930  
Step 4500, loss = -0.1993  
Step 4550, loss = -0.1923  
Step 4600, loss = -0.1925  
Step 4650, loss = -0.1921  
Step 4700, loss = -0.2012  
Step 4750, loss = -0.1943  
Step 4800, loss = -0.1942  
Step 4850, loss = -0.1917  
Step 4900, loss = -0.1922  
Step 4950, loss = -0.1922  
Step 5000, loss = -0.1950  
Step 5050, loss = -0.1912  
Step 5100, loss = -0.1977  
Step 5150, loss = -0.1912  
Step 5200, loss = -0.1955  
Step 5250, loss = -0.2047  
Step 5300, loss = -0.1982  
Step 5350, loss = -0.1918  
Step 5400, loss = -0.1927  
Step 5450, loss = -0.1947  
Step 5500, loss = -0.1905  
Step 5550, loss = -0.1957  
Step 5600, loss = -0.1915  
Step 5650, loss = -0.1949  
Step 5700, loss = -0.1904  
Step 5750, loss = -0.1903

```
Step 5800, loss = -0.1904
Step 5850, loss = -0.1998
Step 5900, loss = -0.1902
Step 5950, loss = -0.1913
Step 6000, loss = -0.1897
Step 6050, loss = -0.1894
Step 6100, loss = -0.1916
Step 6150, loss = -0.1892
Step 6200, loss = -0.1966
Step 6250, loss = -0.1889
Step 6300, loss = -0.1888
Step 6350, loss = -0.1949
Step 6400, loss = -0.1907
Step 6450, loss = -0.1921
Step 6500, loss = -0.1893
Step 6550, loss = -0.1900
Step 6600, loss = -0.1884
Step 6650, loss = -0.2011
Step 6700, loss = -0.1959
Step 6750, loss = -0.1930
Step 6800, loss = -0.1935
Step 6850, loss = -0.1918
Step 6900, loss = -0.1886
Step 6950, loss = -0.1877
Step 7000, loss = -0.1952
Step 7050, loss = -0.1876
Step 7100, loss = -0.1877
Step 7150, loss = -0.1989
Step 7200, loss = -0.1873
Step 7250, loss = -0.1886
Step 7300, loss = -0.1876
Step 7350, loss = -0.1911
Step 7400, loss = -0.1873
Step 7450, loss = -0.1880
Step 7500, loss = -0.1870
Step 7550, loss = -0.1875
Step 7600, loss = -0.1904
Step 7650, loss = -0.1907
Step 7700, loss = -0.1941
Step 7750, loss = -0.1884
Step 7800, loss = -0.1869
Step 7850, loss = -0.1884
Step 7900, loss = -0.1865
Step 7950, loss = -0.1908
```

Our reference solution produces, but don't worry if yours is not exactly the same.

```
Full batch: accuracy: 0.9240, f1_score: 0.9384
```

```
Mini-batch: accuracy: 0.9415, f1_score: 0.9533
```

```
In [12]: y_pred_full = predict(X_test, w_full)
y_pred_minibatch = predict(X_test, w_minibatch)

print('Full batch: accuracy: {:.4f}, f1_score: {:.4f}'
      .format(accuracy_score(y_test, y_pred_full), f1_score(y_test, y_pred_full)))
print('Mini-batch: accuracy: {:.4f}, f1_score: {:.4f}'
      .format(accuracy_score(y_test, y_pred_minibatch), f1_score(y_test, y_pred_minibatch)))
```

```
Full batch: accuracy: 0.9240, f1_score: 0.9384
```

```
Mini-batch: accuracy: 0.9415, f1_score: 0.9533
```

```
In [13]: plt.figure(figsize=[15, 10])
plt.plot(trace_full, label='Full batch')
plt.plot(trace_minibatch, label='Mini-batch')
plt.xlabel('Iterations * 50')
plt.ylabel('Loss  $\mathcal{L}(\mathbf{w})$ ')
plt.legend()
plt.show()
```

